From Discovery to Verifiable Reuse:

SEA Observatory as a Human-Curated Regional Resource for AI Governance

Supheakmungkol Sarin, Adrian Mak, Gordon Cheung, Anonyo Mitra, Wing-Yee Lau, Edward Tsoi, P.T. Nhean, Xuxin Fan

AI Safety Asia

ABSTRACT

Public-sector AI governance in multilingual settings often fails not because relevant materials are absent, but because they are difficult to find, normalize, verify, and reuse under administrative time pressure. This paper reports SEA Observatory, a human-curated regional AI governance resource for Southeast Asia, with three current access surfaces: a public country-level browse surface, a citation-visible multilingual assistant in beta, and GovSim, a standalone scenario-rehearsal module. Its laws-and-policies core draws on official, publicly available materials, including laws, regulations, regulator notices, standards, procurement guidance, strategies, and consultation drafts, and retains a smaller labeled contextual layer. As of May 2026, the resource spans more than 350 indexed materials across 11 Southeast Asian jurisdictions plus ASEAN-level bodies. Items are reviewed before ingestion and normalized against a metadata schema that records source class, jurisdiction, document type, date, language, and source URL, with issuer, document-status labels, and canonical URL retained where available and meaningful. We contribute a practice-informed curation and metadata approach and an implementation account of SEA Observatory's current browse, assistant, and scenario-rehearsal surfaces. The contribution is source-checkable reuse infrastructure rather than a new retrieval model, automated legal-interpretation system, impact evaluation, or end-to-end workflow.

CCS CONCEPTS

• Information systems • Digital government • Information retrieval • Computing methodologies • Natural language processing

KEYWORDS

AI governance, digital government, public records, multilingual retrieval, citation-first AI, system demonstration, policy materials, reuse, scenario rehearsal, Southeast Asia

1 Introduction

A procurement, oversight, or legislative team asks a mundane but consequential question: which current AI governance instrument applies, is it binding or merely consultative, and what source should be cited in a brief, note, or coordination memo? In multilingual settings the documents often exist; what fails is the path from question to reusable answer. Teams still spend time locating the operative record, checking whether it is current, translating it into usable form, distinguishing law from draft or strategy, and reconstructing the source trail by hand.

That bottleneck matters because governance work is accelerating while record use remains fragmented. The 2026 AI Index describes a widening gap between what AI systems can do and how prepared institutions are to govern and evaluate them responsibly [2]. Regional reporting characterizes Southeast Asia as moving quickly from experimentation to deployment [1], while the International AI Safety Report 2026 emphasizes that consequential policy decisions must be made under rapid technical change and incomplete measurement [3]. Multilingual governance work, in that setting, is often not information-poor but workflow-poor. Open-government scholarship makes a related point: publication alone does not guarantee reuse [5].

This paper reports SEA Observatory, a human-curated regional resource for AI governance built for that failure mode. The contribution is neither a new retrieval model nor a completed workflow deployment. It is practice-grounded infrastructure: a reviewed, primary-source-led corpus supplemented by clearly labeled contextual materials, plus access surfaces that keep provenance visible at the point of use. As of May 2026, the current implementation comprises a public country-level browse surface, a citation-visible multilingual assistant in beta, and GovSim, a standalone scenario-rehearsal module. The longer-term path from source discovery to drafting and rehearsal still matters, but it is deliberately separated from the current implementation boundary.

Southeast Asia is the motivating setting because it sharply concentrates the broader governance bottleneck: AI uptake is rising, repositories are fragmented across agencies, languages are numerous, and administrative capacity is uneven. The substantive domain is AI governance, but the underlying usability problem is broader—authoritative public records are hard to find, verify, compare, and safely reuse under time pressure.

2 Resource construction and stakeholder-informed curation

The resource is built around three recurring frictions surfaced in stakeholder consultation and beta cycles. Search friction arises when relevant materials are dispersed across ministries, regulators, standards bodies, gazettes, parliamentary portals, and consultation sites. Language-mediation friction arises when cross-jurisdiction work depends on ad hoc English summaries rather than direct engagement with the source record. Source-type friction arises when a fluent answer does not reveal whether it rests on binding law, draft consultation text, strategy, standard, advisory guidance, or commentary. The resource design drew on AISA regional engagement, practitioner conversations, and beta feedback. Across these inputs, a recurring design requirement was source visibility—issuer, date, status label where available, and URL—at the point of reuse.

The resource is anchored in a laws-and-policies core that is primary-source-led rather than primary-source-only. It prioritizes official, publicly available materials directly relevant to AI governance—laws, regulations, procurement guidance, regulator notices, standards, strategies, and consultation drafts—and retains a smaller set of clearly labeled contextual materials (commentary and think-tank analysis) when these help users navigate a fragmented regional landscape. To mitigate selection bias, contextual inclusions are restricted to materials issued by recognized multilateral bodies or those frequently cited within primary government documents. Items are reviewed for source validity, relevance, duplication risk, and URL accessibility before ingestion, then normalized against a common schema. The schema records source class, jurisdiction, document type, publication or effective date, language, and source URL; primary-source materials additionally preserve issuer, canonical URL where available, and document-status labels retained where available and meaningful. These labels support administrative triage and should not be treated as legal determinations.

In its current form, the resource spans more than 350 indexed materials across 11 Southeast Asian jurisdictions plus ASEAN-level bodies. It is centered on official laws, policies, and related primary sources, with a smaller layer of clearly labeled contextual materials. Multiple regional interface languages are supported on the browse and assistant surfaces, and passage-level segments are retained so generated answers can link back to inspectable evidence rather than only to document titles. The implementation boundary is explicit. Three access surfaces are live in the current implementation: a public country-level browse surface, a citation-visible multilingual assistant in beta, and GovSim, a standalone scenario-rehearsal module. Corpus-wide filtering, query-time restriction to user-selected subsets, direct citation carry-forward into drafting, status-preserving comparison, and end-to-end linkage between retrieval and rehearsal remain part of a later integration path.

3 Current access surfaces

The current implementation exposes the resource through three surfaces: a public country-level browse surface, a citation-visible assistant in beta, and GovSim, a standalone scenario-rehearsal module. They support discovery, question answering, and structured discussion.

Figure 1: SEA Observatory architecture and implementation boundary at submission time. The current resource combines a human-reviewed laws-and-policies core, a smaller clearly labeled contextual layer, and three live access surfaces on the same corpus: a public country-level browse surface, a citation-visible assistant, and a standalone scenario module. Corpus-wide metadata filtering, subset-scoped querying, status-preserving comparison, citation carry-forward, and end-to-end linkage remain part of a later integration path.

3.1 Resource architecture and metadata contract

Query answering operates over the curated corpus rather than the open web. It follows a retrieval-augmented pattern [4], but the important constraint is evidentiary rather than algorithmic: items remain tethered to source-level metadata and passage-level evidence. The metadata contract is operational, not merely descriptive. All ingested items carry source class, jurisdiction, document type, publication or effective date, language, and source URL; primary-source materials additionally retain issuer, document status, and canonical URL where meaningful. Those fields remain attached to passage-level segments so downstream retrieval does not orphan a claim from its source. The multilingual layer draws on established cross-language retrieval practice rather than English-only mediation [7,8].

3.2 Public country-level browse surface

The workflow begins in the country-level observatory view. The user starts with a broad governance issue—vendor disclosure, human-oversight requirements, incident escalation—and narrows from topic to a bounded record set by filtering on jurisdiction, issuer, document type, language, and status. This matters because many public-sector tasks fail before interpretation: a team may know that a jurisdiction has issued relevant guidance and still spend disproportionate time locating the operative version and confirming whether it is official, draft, or merely advisory.

3.3 Citation-backed multilingual assistant

The assistant, currently in beta, lets users query the corpus in natural language and receive citation-backed answers. Questions can be asked in English or other supported languages, and responses rendered in multiple regional languages through the interface or by direct instruction in chat. Multilingual support is provided through LLM-based interpretation of user prompts and answer generation or translation, while citations remain anchored to the original source documents and passages. Each answer is accompanied by source cards showing source class, document type, status where applicable, date, and linked passages. Contextual materials are flagged as such rather than presented as governing instruments. The design target is time-to-verifiable answer, not answer fluency alone.

3.4 GovSim: standalone scenario-rehearsal module

GovSim, the standalone scenario-rehearsal module, is the third surface. Users open curated governance scenarios and structured discussion prompts on issues such as vendor disclosure, escalation paths, infrastructure regulation, and review triggers. The module is intended for scenario learning, policy-drafting coordination, and facilitated discussion rather than automated decision-making; framework references and required deliverables remain visible inside the exercise. In the current implementation, the module operates as a standalone surface rather than as an integrated step in the browse-or-chat workflow. Direct handoff from browse or chat into the scenario view is not yet implemented: users cannot pass a selected evidence packet into the exercise, and record-level citations are not preserved end-to-end.

4 Infrastructure contribution and positioning

Adjacent systems clarify both the fit and the limits of the claim. Baron et al. study machine-learning support for review of one class of FOIA-exempt material, the deliberative process privilege [9]; Governor improves access to open-data portal tables while preserving the provenance of integrated tables [10]; and RequestAtlas supports journalists navigating iterative public-records requests [11]. The nearest comparison class, however, is AI observatories. OECD.AI offers a policy navigator and living repository of policies and initiatives [12]; UNESCO's Global AI Ethics and Governance Observatory foregrounds country readiness information and governance resources [13]; the European Commission's AI Watch monitors AI development, uptake, impact, and policy initiatives in Europe [14]. These are valuable neighboring infrastructures, but their primary unit of value is monitoring, inventory, or readiness. The resource described here is designed instead for source-sensitive governance reuse—preserving source class, provenance, and, where applicable, document status so that materials can be found, queried, inspected, and reused.

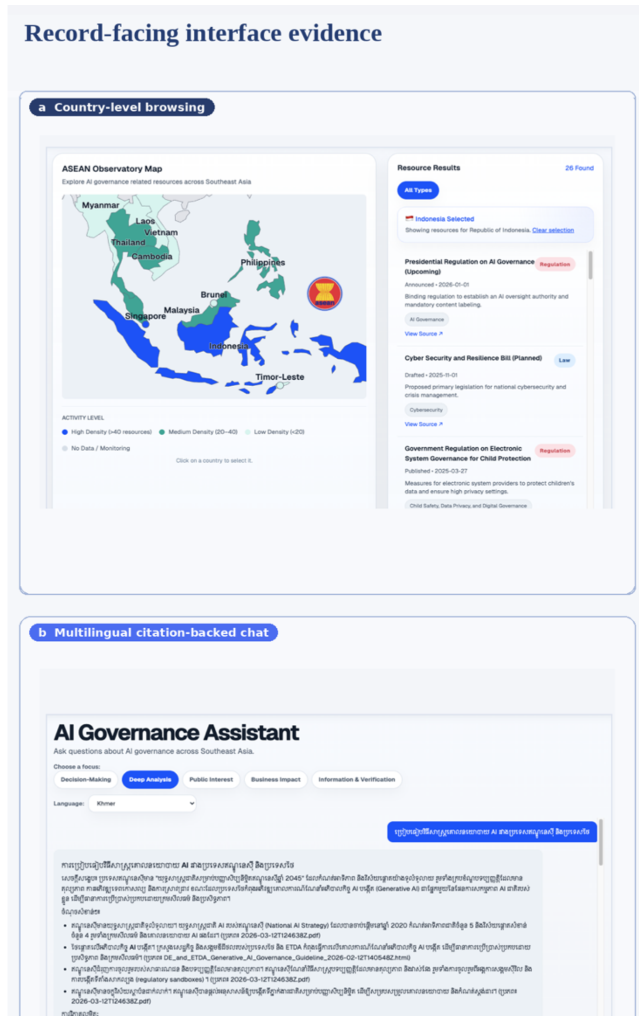

Figure 2: SEA Observatory record-facing surfaces built on top of the regional resource. Panel (a) shows the public country-level browse surface with attached metadata after country selection; panel (b) shows the citation-visible multilingual assistant in beta.

That distinction matters because administrative reuse is where failure becomes costly. A source-sensitive resource is needed not only by national AI offices or regulators, but also by procurement teams, legislative staff, journalists, researchers, and civil-society actors who must inspect or contest governance claims. Multilingual design also matters for whose knowledge stays visible: when regional governance work defaults to English commentary, smaller jurisdictions and non-English official materials are more likely to be underweighted. The novelty claim is therefore infrastructure-level rather than algorithmic. We do not claim a new retrieval model, nor that the present implementation is already proven superior to neighboring systems. We argue instead that a regional AI governance resource becomes more usable for source-sensitive reuse when its corpus is human-curated, primary-source-led, and transparently labeled, and when its access surfaces preserve issuer, date, status (where applicable), and source linkage at the point of use.

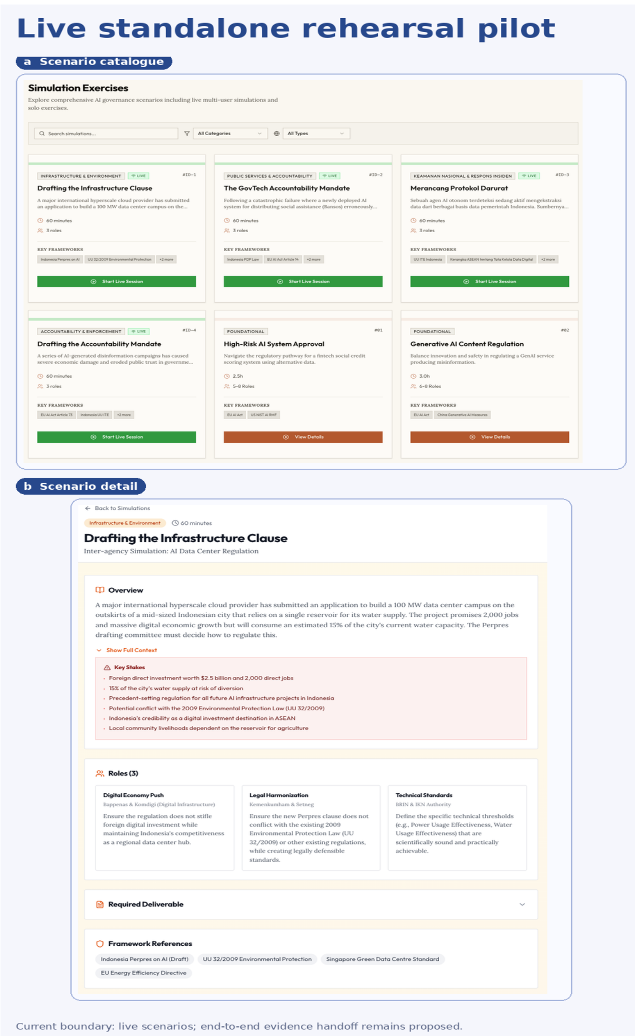

Figure 3: GovSim, the standalone scenario-rehearsal module within SEA Observatory. Panel (a) shows the scenario catalogue; panel (b) shows a representative scenario detail view with roles, stakes, deliverable requirements, and framework references. The module is live as a separate surface, but end-to-end evidence handoff from record selection or chat into the exercise is not yet implemented.

5 Limits

The current implementation has three limits. First, the scenario-rehearsal module is separate: evidence packets are not handed off automatically from browse or chat, and citations do not persist throughout the exercise. Second, coverage depends on materials that can be collected, reviewed, and normalized reliably from public sources. Third, retrieval and translation errors remain possible, especially for legally dense or jurisdiction-specific materials. Visible citations reduce but do not eliminate interpretive risk; human-in-the-loop review remains essential before reuse in a memo, procurement note, legislative comparison, or oversight question [6]. SEA Observatory should therefore be read as source-checkable governance infrastructure, not as legal advice, an impact evaluation, or an integrated end-to-end workflow.

6 Conclusion

Fast-moving governance settings need more than access to documents. They need regional resources that turn fragmented policy materials into evidence that can be found, checked, and reused under administrative time pressure.

This paper reported SEA Observatory, a human-curated regional AI governance resource for multilingual use, together with a public country-level browse surface, a citation-visible assistant in beta, and a standalone scenario-rehearsal module. The current contribution is bounded: curation, metadata design, and implementation description. It does not claim measured improvement, automated legal interpretation, policy impact, or an integrated workflow. Further work is needed to test retrieval quality, citation faithfulness, workflow efficiency, and institutional uptake. Looking ahead, pairing SEA Observatory's curated record with GovSim's structured scenario rehearsal points toward a complementary mode of regional governance work—supporting both source-level discovery of binding instruments and the cross-team collaboration that turns those instruments into shared practice.

REFERENCES

[1] Vinayak HV, Vivek Lath, Amy Yu, and Saurish Basu. 2026. AI in Southeast Asia: An Era of Opportunity. McKinsey & Company, in collaboration with Singapore Economic Development Board and Tech in Asia.

[2] Sha Sajadieh, Loredana Fattorini, Raymond Perrault, Yolanda Gil, Vanessa Parli, Lapo Santarlasci, Juan Pava, Nestor Maslej, Russ Altman, Erik Brynjolfsson, Carla Brodley, Jack Clark, Virginia Dignum, Vipin Kumar, James Landay, Terah Lyons, James Manyika, Juan Carlos Niebles, Yoav Shoham, Elham Tabassi, Russell Wald, Toby Walsh, and Dan Weld. 2026. The AI Index 2026 Annual Report. AI Index Steering Committee, Institute for Human-Centered AI, Stanford University.

[3] Yoshua Bengio et al. 2026. International AI Safety Report 2026. UK Department for Science, Innovation and Technology.

[4] Patrick Lewis, Ethan Perez, Aleksandara Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen-tau Yih, Tim Rocktaschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In Advances in Neural Information Processing Systems 33, 9459–9474.

[5] Marijn Janssen, Yannis Charalabidis, and Anneke Zuiderwijk. 2012. Benefits, Adoption Barriers and Myths of Open Data and Open Government. Information Systems Management 29(4), 258–268.

[6] National Institute of Standards and Technology. 2023. Artificial Intelligence Risk Management Framework (AI RMF 1.0). U.S. Department of Commerce.

[7] Luiz Bonifacio, Vitor Jeronymo, Hugo Queiroz Abonizio, Israel Campiotti, Marzieh Fadaee, Roberto Lotufo, and Rodrigo Nogueira. 2021. mMARCO: A Multilingual Version of the MS MARCO Passage Ranking Dataset. arXiv:2108.13897.

[8] Xinyu Zhang, Nandan Thakur, Odunayo Ogundepo, Ehsan Kamalloo, David Alfonso-Hermelo, Xiaoguang Li, Qun Liu, Mehdi Rezagholizadeh, and Jimmy Lin. 2022. Making a MIRACL: Multilingual Information Retrieval across a Continuum of Languages. arXiv:2210.09984.

[9] Jason R. Baron, Mahmoud F. Sayed, and Douglas W. Oard. 2022. Providing More Efficient Access to Government Records: A Use Case Involving Application of Machine Learning to Improve FOIA Review for the Deliberative Process Privilege. ACM Journal on Computing and Cultural Heritage 15, 1 (2022), 1-19.

[10] Chang Liu, Arif Usta, Jian Zhao, and Semih Salihoglu. 2023. Governor: Turning Open Government Data Portals into Interactive Databases. In Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems, 415:1-415:16.

[11] Rachel B. Warren, Lisa Pickoff-White, Aditya G. Parameswaran, and Niloufar Salehi. 2025. RequestAtlas: Supporting the Slow and Iterative Process of Requesting Public Records. Proc. ACM Hum.-Comput. Interact. 9, 2 (2025), 1-35.

[12] OECD. 2026. The OECD.AI Policy Navigator. OECD.AI. https://oecd.ai/en/dashboards/policy-initiatives.

[13] UNESCO. 2026. Global AI Ethics and Governance Observatory. UNESCO. https://www.unesco.org/ethics-ai/en.

[14] European Commission Joint Research Centre. 2026. AI Watch. European Commission. https://ai-watch.ec.europa.eu/index_en.